

# Hotel Booking Demand Clustering Study

A Reproducible Experimental Study of Clustering

Alexandre Santos (72970)

David Natal (72997)

Miguel Mestre (73018)

5 June 2026

Unsupervised Learning – 2025/2026

Repository tag: Submission

Repository: <https://github.com/alexandreffs/UnsupervisedLearningProj>

## Abstract

This report presents a reproducible clustering study of hotel booking demand data with the goal of partitioning bookings into interpretable guest segments based solely on pre-stay observable features. Two representations were studied: R-EUCLID-robust-weighted-withADR and a pipeline that excludes the daily rate R-EUCLID-robust-weighted-noADR. K-Means ( $k = 3$ , Silhouette = 0.452, ARI = 0.989 across 10 seeds), iK-Means (automatic  $k = 8$ ), and Agglomerative Hierarchical clustering (Ward, Silhouette = 0.453 at  $k = 3$ ) are compared. The selected three-cluster solution identifies a Non-Refundable Advance Booker segment (12.4%), a Standard Flexible City Guest segment (70.7%), and a Premium Resort Long-Stay segment (16.8%). The noADR comparison confirms that the commitment axis (`is_non_refund`) drives cluster separation independently of price.

## Contents

|          |                                       |          |
|----------|---------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                   | <b>2</b> |
| 1.1      | Context                               | 2        |
| 1.2      | Segmentation question                 | 2        |
| 1.3      | Why clustering is appropriate         | 2        |
| 1.4      | Research questions                    | 3        |
| <b>2</b> | <b>Dataset and Preprocessing</b>      | <b>3</b> |
| 2.1      | Dataset documentation                 | 3        |
| 2.2      | Data quality                          | 3        |
| 2.3      | Leakage control                       | 4        |
| 2.4      | Feature engineering                   | 4        |
| 2.5      | Representations and distance          | 4        |
| <b>3</b> | <b>Clustering Methods</b>             | <b>4</b> |
| 3.1      | K-Means baseline                      | 4        |
| 3.2      | iK-Means                              | 5        |
| 3.3      | Agglomerative hierarchical clustering | 5        |
| <b>4</b> | <b>Experimental Protocol</b>          | <b>5</b> |
| 4.1      | Internal validity                     | 5        |
| 4.2      | Family-specific diagnostics           | 5        |
| 4.3      | Stability and sensitivity             | 5        |
| 4.4      | Reproducibility                       | 5        |
| <b>5</b> | <b>Results</b>                        | <b>5</b> |
| 5.1      | Data preparation outcomes             | 5        |
| 5.2      | Model comparison                      | 6        |
| 5.3      | Cluster profiles                      | 6        |
| 5.4      | Cluster interpretation                | 6        |
| 5.5      | noADR comparison                      | 7        |

|          |                                                      |          |
|----------|------------------------------------------------------|----------|
| <b>6</b> | <b>Discussion</b>                                    | <b>7</b> |
| 6.1      | Which result best answers the segmentation question? | 7        |
| 6.2      | Influence of preprocessing                           | 7        |
| 6.3      | Stability and sensitivity                            | 7        |
| 6.4      | Decision-relevant recommendations                    | 7        |
| 6.5      | What should not be concluded                         | 8        |
| <b>7</b> | <b>Limitations, Ethics, and Reproducibility</b>      | <b>8</b> |
| 7.1      | Limitations                                          | 8        |
| 7.2      | Ethics and governance                                | 8        |
| 7.3      | Reproducibility checklist                            | 8        |
| 7.4      | Academic integrity disclosure                        | 8        |
| <b>8</b> | <b>Conclusion</b>                                    | <b>8</b> |
| <b>A</b> | <b>Feature Table</b>                                 | <b>9</b> |
| <b>B</b> | <b>K-Means Stability Summary</b>                     | <b>9</b> |
| <b>C</b> | <b>Scaler Sensitivity</b>                            | <b>9</b> |

# 1 Introduction

## 1.1 Context

Hotel revenue management depends on understanding who books, when they book, and how their bookings differ in value and service requirements. Booking records contain pre-arrival signals such as rate, stay duration, booking horizon, deposit commitment, meal plan, parking need, and special requests. These signals can reveal latent booking types that support pricing, upselling, and operational planning decisions.

## 1.2 Segmentation question

The concrete segmentation question addressed in this project is:

Can hotel booking records be partitioned into distinct and interpretable behavioural segments using only pre-stay observable characteristics, such as daily rate, total estimated spend, length of stay, booking horizon, meal plan, deposit commitment, special requests, and parking need, in order to identify guest profiles that support targeted revenue-management and personalised-service strategies?

The unit of analysis is the individual booking record. A guest may appear more than once in the dataset, so the clusters should be interpreted as booking segments rather than permanent guest identities.

## 1.3 Why clustering is appropriate

In this project, we don't know beforehand which group each customer belongs to. Our goal is to discover these groups rather than making predictions. Therefore, unsupervised techniques, like clustering, are the right choice.

Clustering is ideal here because hotel bookings have very different characteristics from one another: some people book far in advance, others the day before, there are differences in prices, length of stay, type of deposit paid, and special requests.

## 1.4 Research questions

The report addresses the course research questions as follows:

- **RQ1 – Problem framing:** define a booking-level segmentation question and justify clustering.
- **RQ2 – Representation:** define leakage-safe preprocessing and feature representations for mixed-type booking data.
- **RQ3 – Modelling:** compare K-Means, iK-Means, and Ward hierarchical clustering.
- **RQ4 – Interpretation:** profile the discovered clusters as booking segments.
- **RQ5 – Robustness:** evaluate sensitivity to seeds, subsampling, scaling, feature weighting, and representation changes.

## 2 Dataset and Preprocessing

### 2.1 Dataset documentation

Table 1: Dataset documentation and traceability.

| Item                | Details                                                                                  |
|---------------------|------------------------------------------------------------------------------------------|
| Dataset             | Hotel Booking Demand                                                                     |
| Version             | Course release v1                                                                        |
| Unit of analysis    | One row corresponds to one booking record                                                |
| Original size       | 119,390 rows and 32 columns                                                              |
| Time span           | Arrivals from July 2015 to August 2017                                                   |
| Source              | António, de Almeida, and Nunes (2019); dataset provided through the course platform      |
| Raw-data governance | Raw CSV is not committed to the repository; it must be placed in <code>data/</code>      |
| Checksum            | SHA-256: 7c2ae42a7353905ea136e5c2287f17c92c5435826598bfbb8491c6f0c7b1fc06                |
| Usage terms         | Dataset used under the course release terms, no raw data redistributed in the repository |

### 2.2 Data quality

The main data-quality issues identified are:

- `children` has 4 missing values, imputed as 0 because missing children counts are interpreted as no children in the booking.
- `country` has 488 missing values and is retained only for post-hoc profiling, not as a clustering input.
- `agent` and `company` have high missingness and high cardinality; both are dropped as ID-like fields.
- `adr` has a heavy right tail (maximum approximately €5,400) and one negative value, requiring robust preprocessing and careful interpretation.
- Rows with `adr` = 0 are excluded from clustering because they correspond to complementary or test-like bookings. This removes 1,960 records, leaving 117,430 analytic records.

Table 2: Main missingness and governance decisions.

| Variable              | Missing | Role                | Decision                                      |
|-----------------------|---------|---------------------|-----------------------------------------------|
| <code>children</code> | 4       | Numerical           | Impute as 0                                   |
| <code>country</code>  | 488     | Categorical         | Exclude from clustering; retain for profiling |
| <code>agent</code>    | 16,340  | ID-like categorical | Drop due to high missingness/cardinality      |
| <code>company</code>  | 112,593 | ID-like categorical | Drop due to high missingness/cardinality      |

## 2.3 Leakage control

The segmentation index time is the booking creation time. Therefore, only variables available at or before booking time can be used to form clusters. The following variables are excluded from clustering inputs:

- `is_canceled`: outcome variable.
- `reservation_status`: post-decision/post-arrival status.
- `reservation_status_date`: date associated with the realised status.
- `agent` and `company`: high-cardinality ID-like fields.
- `country`: excluded from clustering to avoid geographical profiling effects; retained only for post-hoc descriptive analysis.

## 2.4 Feature engineering

The final clustering features are:

- `total_nights` = `stays_in_week_nights` + `stays_in_weekend_nights`.
- `total_spend_log` =  $\log(1 + \text{adr}) \times \text{total\_nights}$ , used as a compressed estimated-spend proxy.
- `is_non_refund` = 1 if `deposit_type` is Non Refund, 0 otherwise.
- `lead_time_rank`, an ordinal quantile binning of `lead_time`: 0 for last-minute, 1 for mid-horizon, and 2 for advance bookings.
- `meal_encoded`, an ordinal encoding of meal-plan inclusivity: SC/Undefined = 0, BB = 1, HB = 2, FB = 3.
- `total_of_special_requests`.
- `required_car_parking_spaces`.

## 2.5 Representations and distance

Two Euclidean representations are studied:

Table 3: Final representations.

| Representation                   | Definition                                                                                                                                                                                                                                                                                                                                                                                                 |
|----------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| R-EUCLID-robust-weighted-withADR | <b>Features:</b> <code>total_spend_log</code> , <code>total_nights</code> , <code>is_non_refund</code> , <code>total_of_special_requests</code> , <code>required_car_parking_spaces</code> , <code>meal_encoded</code> , <code>lead_time_rank</code> .<br><b>Scaling/weights:</b> RobustScaler on continuous variables; <code>is_non_refund</code> $\times 3.0$ ; <code>meal_encoded</code> $\times 0.5$ . |
| R-EUCLID-robust-weighted-noADR   | <b>Features:</b> same as R-EUCLID-robust-weighted-withADR, excluding <code>total_spend_log</code> .<br><b>Scaling/weights:</b> same scaling and weights as R-EUCLID-robust-weighted-withADR.                                                                                                                                                                                                               |

All models and internal validity indices are computed in the same weighted Euclidean space used to fit the corresponding clustering model.

# 3 Clustering Methods

## 3.1 K-Means baseline

K-Means is used as the mandatory baseline. The predefined search interval is  $k \in \{2, \dots, 10\}$ , chosen because  $k = 2$  is the minimum meaningful segmentation and values above 10 are unlikely to remain managerially interpretable. The configuration is `n_init=10`, `max_iter=300`, and `random_state=42` for the main run (Runtime:  $\sim 5$  seconds).

## 3.2 iK-Means

Intelligent K-Means (iK-Means) was included as an exploratory method to automatically discover cluster centers. Using Mirkin’s procedure with `min_cluster_size=1000`, the algorithm automatically extracted  $k = 8$  clusters (Runtime:  $\sim 3$  seconds). However, iK-Means was not selected as the primary segmentation model because 8 clusters overly fragment the main booking profiles, adding complexity without offering qualitatively new, actionable insights for revenue management.

## 3.3 Agglomerative hierarchical clustering

Agglomerative hierarchical clustering with Ward linkage is used as the alternative clustering family. Ward linkage is compatible with the chosen representation because the features are embedded in a Euclidean vector space. Due to computational cost, Ward clustering is run on a reproducible subsample of 30,000 records. The main subsample uses seed 42, and robustness is assessed using additional subsamples.

# 4 Experimental Protocol

## 4.1 Internal validity

Models are evaluated using:

- **Silhouette score** as the primary compactness/separation index.
- **Calinski-Harabasz index** as an additional compactness/separation index, where higher is better.
- **Davies-Bouldin index**, where lower is better.

## 4.2 Family-specific diagnostics

For hierarchical clustering, a Ward dendrogram is used to inspect merge structure and support the choice of the cut. For iK-Means, the automatic number of clusters and any anomalous extracted clusters should be documented.

## 4.3 Stability and sensitivity

The robustness protocol includes:

- K-Means repeated across 10 random seeds for  $k \in \{2, 3, 4, 5\}$ .
- Adjusted Rand Index (ARI) between runs as the stability measure.
- Hierarchical clustering repeated on controlled subsamples.
- RobustScaler versus StandardScaler as a preprocessing sensitivity check.
- withADR versus noADR comparison to test whether clusters are driven only by price.

## 4.4 Reproducibility

All experiments are intended to be reproducible from the repository using:

- `experiments.csv` for experiment monitorability (sample\_rule, representation\_id, model, parameters, runtime, seed, silhouette, calinski, davies, diagnostics, notes).
- `requirements.txt` file for explicit dependency management.
- `run_all.py` as the single-entry point, which explicitly reproduces the entire data mining pipeline (preprocessing, modelling, logging, and figures) from scratch.

# 5 Results

## 5.1 Data preparation outcomes

After excluding zero-ADR records, the analytic dataset contains 117,430 bookings. The withADR matrix has 7 features and the noADR matrix has 6 features.

## 5.2 Model comparison

Table 4: Main model comparison: quantitative results.

| Method   | Setting                         | Sil.  | CH     | DB    | Stability              |
|----------|---------------------------------|-------|--------|-------|------------------------|
| K-Means  | $k = 3$ , withADR               | 0.452 | 63,837 | 0.857 | ARI $0.989 \pm 0.007$  |
| K-Means  | $k = 2$ , withADR               | 0.392 | 54,947 | 1.074 | ARI $0.998 \pm 0.002$  |
| K-Means  | $k = 3$ , noADR                 | 0.443 | 60,587 | 0.928 | n/a                    |
| iK-Means | $k = 8$ (auto)                  | 0.327 | 56,165 | 1.012 | n/a                    |
| Ward     | $k = 3$ , withADR, $n = 30,000$ | 0.453 | 15,928 | 0.832 | Sil. $0.445 \pm 0.014$ |

Table 5: Main model comparison: interpretation.

| Model / setting            | Observation                                                                                                                                                                                                               |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| K-Means, $k = 3$ , withADR | Selected solution. It gives the best balance between internal validity, seed stability, and business interpretability.                                                                                                    |
| K-Means, $k = 2$ , withADR | Stable but too coarse: it merges the Premium Resort Long-Stay profile with the Standard Flexible City Guest profile and loses useful granularity.                                                                         |
| K-Means, $k = 3$ , noADR   | Useful sensitivity check. The noADR solution separates a Special Requests group, a Standard Flexible group, and a Non-Refundable high-lead-time group, confirming that the commitment axis remains visible without price. |
| iK-Means                   | Automatic $k = 8$ fragments the main groups into smaller clusters. It is therefore useful as a diagnostic but is not selected as the primary segmentation because it does not reveal qualitatively new booking profiles.  |
| Ward, $k = 3$ , withADR    | Conceptually distinct method that corroborates the same three-cluster structure on a reproducible subsample.                                                                                                              |

The selected solution is K-Means with  $k = 3$  on the withADR representation. This choice is justified because it achieves the strongest balance between internal validity, stability, and interpretability. Although  $k = 2$  is stable, it merges two managerially different segments. Values above  $k = 3$  tend to fragment the same business profiles without producing clearer narratives. iK-Means is not selected as the primary solution because its automatic  $k = 8$  solution fragments the three main groups without revealing qualitatively new profiles.

## 5.3 Cluster profiles

Table 6: Selected K-Means cluster profiles for  $k = 3$  withADR.

| Cluster | Size   | %    | Main characteristics                                                                                                   | Label                         |
|---------|--------|------|------------------------------------------------------------------------------------------------------------------------|-------------------------------|
| C2      | 14,513 | 12.4 | ADR approximately €90, <code>is_non_refund=1.0</code> , stay 2.7 nights, advance lead-time profile                     | Non-Refundable Advance Booker |
| C1      | 83,139 | 70.7 | ADR approximately €103, flexible deposit, stay 2.6 nights, last-minute/mid lead-time profile, highest parking need     | Standard Flexible City Guest  |
| C0      | 19,778 | 16.8 | ADR approximately €114, stay 7.7 nights, total spend 136% above the mean, most special requests, resort-hotel dominant | Premium Resort Long-Stay      |

## 5.4 Cluster interpretation

The **Non-Refundable Advance Booker** segment is defined by full deposit commitment (`is_non_refund=1.0` for all members), a more advanced lead\_time, and short stays of about 2.7 nights.

The **Standard Flexible City Guest** segment forms the majority of bookings (70.7%). These records are typically flexible, short-stay bookings with no non-refundable commitment, an average stay of about 2.6 nights, and the highest parking-need proportion. This segment likely contains both business and urban leisure demand.

The **Premium Resort Long-Stay** segment is the smallest but most valuable profile. It is characterised by longer stays of about 7.7 nights, ADR around €114, total estimated spend 136% above the grand mean, and more special requests. This segment is suitable for pre-arrival personalisation and upselling.

## 5.5 noADR comparison

The noADR representation tests whether clusters are merely price bands. Removing `total_spend_log` reduces the explicit price/value signal. The Non-Refundable segment remains structurally meaningful, indicating that deposit commitment is an independent axis of segmentation. However, separation between Standard and Premium bookings becomes weaker because price and total spend are important for distinguishing high-value long-stay behaviour.

# 6 Discussion

## 6.1 Which result best answers the segmentation question?

K-Means with  $k = 3$  on the withADR representation best answers the segmentation question. It provides three stable and interpretable booking profiles that align with actionable revenue-management decisions: commitment-based rate locking, flexible city-hotel demand management, and high-value resort-stay personalisation.

## 6.2 Influence of preprocessing

The feature weight on `is_non_refund` is one of the most consequential preprocessing choices. Since it is a binary variable, its natural range is narrower than the continuous features; however, deposit commitment is assumed to be a key axis for differentiating booking tiers and should not be underweighted relative to price or stay duration. A multiplier of 3.0 was therefore applied, with the specific value determined by experimentation. The meal-plan weight of 0.5 reduces the influence of an ordinal encoding that, while useful for ordering, is less decisive for defining pricing tiers. These weights make the Euclidean metric more aligned with the business segmentation question, but they also introduce subjective assumptions.

## 6.3 Stability and sensitivity

K-Means is highly stable across seeds for the selected  $k = 3$  solution. Hierarchical Ward clustering provides independent support for the three-cluster structure. Sensitivity checks using scaler changes and the noADR representation indicate that the main qualitative conclusions are not driven solely by a single preprocessing decision, although the distinction between Standard and Premium segments depends partly on value-related features.

## 6.4 Decision-relevant recommendations

- **Non-Refundable Advance Booker:** offer early-bird packages and prepaid upgrades at booking time.
- **Standard Flexible City Guest:** use targeted flexibility-based promotions, parking-inclusive packages, and last-minute offers.
- **Premium Resort Long-Stay:** prioritise pre-arrival personalisation, room upgrades, spa/F&B offers, and long-stay service planning.

## 6.5 What should not be concluded

The clusters are not externally validated customer types. They are booking-level segments, not permanent guest identities. A returning guest may appear in different clusters across different stays. Internal validity indices measure geometric compactness and separation, not business profitability or causal effects.

## 7 Limitations, Ethics, and Reproducibility

### 7.1 Limitations

- There are no ground-truth segment labels.
- The unit of analysis is a booking record, not a unique guest.
- The meal-plan encoding assumes an ordinal structure and equal spacing between meal levels.
- Lead-time quantile binning simplifies a highly skewed continuous variable.
- Manual feature weights encode subjective modelling assumptions.
- The dataset concerns two Portuguese hotels from 2015–2017, limiting generalisability.
- iK-Means produced an automatic  $k = 8$  solution, but this result should be treated as exploratory because it fragments the main segments.

### 7.2 Ethics and governance

The dataset is used according to the course release terms. Raw data are not committed to the repository. No direct personal identifiers are used.

### 7.3 Reproducibility checklist

The final repository contains:

- `requirements.txt` for explicit dependency management.
- `run_all.py` as the single-entry point to execute the pipeline from scratch.
- `experiments.csv` as the experiment log, recording methods, parameters, runtimes, and metrics.
- Figures and tables are generated dynamically and displayed directly within the Jupyter Notebook upon execution.

### 7.4 Academic integrity disclosure

Generative-AI tools were consulted only for methodological discussion, syntax review, and report-structure support.

## 8 Conclusion

This study groups 117,430 hotel bookings based on information available before the guests' stay. Our main model, K-Means with 3 clusters, successfully identified three clear guest profiles: Non-Refundable Advance Booker, Standard Flexible City Guest, and Premium Resort Long-Stay. These results are reliable, as confirmed by our internal metrics, stability tests, and comparison with a hierarchical clustering method. Comparing models with and without the room price (ADR) showed that guest commitment (such as lead time and deposits) is always a key factor. However, including the price provides better insights for revenue management. As a next step, we considered that this analysis should be applied to broader datasets, including more hotels and wider date ranges, to confirm if these guest profiles remain consistent.



## A Feature Table

Table 7: withADR feature table.

| Feature                     | Type         | Scaling      | Weight       |
|-----------------------------|--------------|--------------|--------------|
| total_spend_log             | Continuous   | RobustScaler | $\times 1.0$ |
| total_nights                | Continuous   | RobustScaler | $\times 1.0$ |
| is_non_refund               | Binary       | Raw          | $\times 3.0$ |
| total_of_special_requests   | Count        | Raw          | $\times 1.0$ |
| required_car_parking_spaces | Binary/count | Raw          | $\times 1.0$ |
| meal_encoded                | Ordinal      | Raw          | $\times 0.5$ |
| lead_time_rank              | Ordinal      | Raw          | $\times 1.0$ |

## B K-Means Stability Summary

Table 8: K-Means stability across 10 seeds.

| k | Mean ARI | Std ARI | Mean Sil. | Std Sil. | Notes                                           |
|---|----------|---------|-----------|----------|-------------------------------------------------|
| 2 | 0.9975   | 0.0021  | 0.388     | 0.004    | Stable but under-segmented<br>Selected solution |
| 3 | 0.989    | 0.007   | 0.450     | 0.002    |                                                 |
| 4 | 0.9942   | 0.0057  | 0.375     | 0.003    | Additional fragmentation<br>Less interpretable  |
| 5 | 0.9803   | 0.0138  | 0.327     | 0.002    |                                                 |

## C Scaler Sensitivity

Table 9: Scaler sensitivity for numerical features,  $k = 3$ .

| Scaler         | Silhouette | CH     | DB    |
|----------------|------------|--------|-------|
| RobustScaler   | 0.450      | 63,837 | 0.857 |
| StandardScaler | 0.443      | 62,062 | 0.906 |

## References

- [1] N. António, A. de Almeida, and L. Nunes. Hotel booking demand datasets. *Data in Brief*, 22:41–49, 2019. <https://doi.org/10.1016/j.dib.2018.11.126>
- [2] B. Mirkin. *Clustering for Data Mining: A Data Recovery Approach*. Chapman & Hall/CRC, 2005.
- [3] L. Kaufman and P. J. Rousseeuw. *Finding Groups in Data: An Introduction to Cluster Analysis*. Wiley, 1990.
- [4] F. Pedregosa et al. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.